

# Homework 2

Christopher Mertin  
CS6966: Theory of Machine Learning

March 8, 2017

6. (Convexity basics) For this problem, let  $f$  be a convex function defined over a convex set  $K$ , and suppose the diameter of  $K$  is 1.

- (a) Let  $x \in K$ , and suppose  $f(x) = 2$  and  $\|\nabla f(x)\| = 1$ . Give a lower bound on  $\min_{z \in K} f(z)$ .

**Solution:**

- (b) Let  $x^*$  be the minimizer of  $f$  over  $K$  (suppose it is unique), and let  $x$  be any other point. The intuition behind gradient descent is that the vector:  $-\nabla f(x)$  points *towards*  $x^*$ . Prove that this is indeed true, in the sense that  $\langle \nabla f(x), x - x^* \rangle > 0$  (*i.e.*, the negative gradient makes an acute angle with the line to the optimum).

**Solution:**

- (c) Suppose now that the function  $f$  is *strictly convex*, *i.e.*,  $f(\lambda x + (1 - \lambda)y) < \lambda f(x) + (1 - \lambda)f(y)$  (strictly), for all  $x \neq y$ , and  $0 < \lambda < 1$ .

Prove that all the *maximizers* of  $f$  over  $K$  lie on the boundary of  $K$ .

[*Hint:* You may want to use the definition that a point  $x$  is not on the boundary iff there exists points  $y, z \in K$  such that  $x = (y + z)/2$ ]

**Solution:**

7. (Gradient Descent Basics)

- (a) Give an example of a function defined over  $\mathbb{R}$ , for which for *any* step-size  $\eta > 0$  (no matter how small), gradient descent with step size  $\eta$  oscillates around the optimum point (*i.e.*, never gets to distance  $< \eta/4$  to it), for some starting point  $x \in \mathbb{R}$ .

**Solution:**

- (b) Consider the function  $f(x, y) = x^2 + y^2/4$ , and suppose we run gradient descent with starting point  $(1, 1)$ , and  $\eta = 1/4$ . Do we get arbitrarily close to the minimum? Experimentally, find the *threshold* for  $\eta$ , beyond which gradient descent starts to oscillate.

**Solution:**

- (c) Why is the behavior similar to that in part (a) (oscillation for *every*  $\eta$ ) not happening in part (b)?

**Solution:**

8. (Stochastic Gradient Descent) Suppose we have points  $\{(a_1, b_1), (a_2, b_2), \dots, (a_n, b_n)\}$  in the plane, and suppose that  $|a_i| \leq 1$ , and  $|b_i| \leq 1$  for all  $i$ . Let  $f(x, y) = \frac{1}{n} \sum_{i=1}^n f_i(x, y)$ , where  $f_i(x, y) = (x - a_i)^2 + (y - b_i)^2$ .

- (a) What is the point  $(x, y)$  that minimizes  $f(x, y)$ ?

**Solution:**

- (b) Suppose we perform gradient descent (on  $f$ ) with step size  $0 < \eta < 1$ . Give a geometric interpretation for one iteration.

**Solution:**

- (c) Now suppose we perform stochastic gradient descent with fixed step size  $0 < \eta < 1$ , and by picking  $i$  at random in  $\{1, 2, \dots, n\}$ , and moving along the gradient of  $f_i$  (as in SGD seen in class). After  $T$  steps, for  $T$  large enough, can we say that we get arbitrarily close to the optimum? (Provide a clear explanation)

[Hint: Remember  $\eta$  is fixed]

**Solution:**

- (d) Pick  $n = 100$  random points in  $[-1, 1]^2$  (uniformly), and run SGD for fixed  $\eta = 1/2$ , as above. Write down what the distance to optimum is, after  $T = \{10, 100, 1000\}$  iterations (if you want to be careful, you should average over 5 random choices for the initialization). Now consider dropping the step size  $\eta_t = 1/t$ , and write down the result for  $T$  as above.

**Solution:**

9. (Numeric accuracy in MW updates) Consider the randomized experts setting we saw in class (we maintain a distribution over experts at each time, and the loss of the algorithm at that time is the expected loss over the distribution). Consider the simple setting where the experts predict 0/1, and the loss is either 0 or 1 for each expert. We saw how to update the probabilities (multiply by  $e^{-\eta}$  if an expert makes a mistake, keep unchanged otherwise, and renormalize). One of the common problems here is that numeric errors in such computations tend to compound if not done carefully.

Suppose we have  $N$  experts, and we start with a uniform distribution over all of them. Let  $p_t^{(i)}$  denote the probability of expert  $i$  at time  $t$ , for the “True” (infinite precision) multiplicative weight algorithm, and let  $q_t^{(i)}$  denote the probabilities that the “real life” algorithm uses (due to precision limitations).

- (a) Consider one simple suggestion: Say we zero out weights that are “too small,” specifically, suppose we set  $q_t^{(i)} = 0$  if  $q_t^{(i)} / \max_j q_t^{(j)} < \epsilon$ , for some precision parameter  $\epsilon$  (such changes frequently occur due to roundoff). Other than this, suppose that the  $q_t^{(i)}$  are updated accurately. Prove that in this case, we cannot hope to achieve any non-trivial regret bound. Specifically, for large enough  $T$ ,

the algorithm can have error  $T(1 - \mathcal{O}(1))$ , while the best expert may have error  $\mathcal{O}(T)$ .

[*Hint*: In this case, we are “losing” all information about an expert]

**Solution:**

- (b) A simple way to avoid this (in this setting) is to avoid storing probabilities, but instead maintaining only the number of mistakes  $m_t^{(i)}$ . Prove how this suffices to recover the probabilities  $p_t^{(i)}$  (assuming infinite precision arithmetic).

**Solution:**

- (c) Suppose we use the idea in part (b) to construct a distribution  $q_t$  that differs from  $p_t$  by  $< \epsilon$  in the  $\ell_1$  norm, i.e.,  $\sum_i |p_t^{(i)} - q_t^{(i)}| < \epsilon$ . Then, assuming we construct such a  $q_t$  at time  $t$  to sample, show that the expected number of mistakes of the algorithm is bounded by  $(1 + \eta) \min_i m_T^{(i)} + \mathcal{O}(\log(N)/\eta) + \epsilon T$ .

**Solution:**

- (d) The bound above is not great if there is an expert who makes very small number of mistakes (compared to  $T$ ). Noting that we are dealing with binary predictions, can you come up with a way to run the algorithm, so as to obtain a mistake bound of  $(1 + \eta + 2\epsilon) \min_i m_T^{(i)} + \mathcal{O}(\log(N)/\eta)$ ?

**Solution:**