

Homework 3

Christopher Mertin
CS6966: Theory of Machine Learning

April 4, 2017

10. Consider the simple experts setting: we have n experts E_1, \dots, E_n , and each one makes a 0/1 prediction each morning. Using these predictions, we need to make a prediction each morning, and at the end of the day we get a loss of 0 if we predicted right, and 1 if we made a mistake. This goes on for T days.

Consider an algorithm that at every step, goes with the prediction of the “best” (*i.e.* the one with the least mistakes so far) expert so far. Suppose that ties are broken by picking the expert with a smaller index. Given an example in which this strategy can be really bad – specifically, the number of mistakes made by the algorithm is roughly a factor of n worse than that of the best expert in hindsight.

Solution:

11. We saw in class a proof that the VC dimension of the class of n -node, m -edge *threshold* neural networks is $\mathcal{O}((m+n)\log(n))$. Let us give a “counting” proof, assuming the weights are binary (0/1). (This is often the power given by VC dimension based proofs – they can “handle” continuous parameters that cause problems for counting arguments).

- (a) Specifically, how many “network layouts” can there be with n nodes and m edges? Show that $\binom{n(n-1)/2}{m}$ is an upper bound.

Solution:

- (b) Given a network layout, argue that the number of “possible networks” is at most $2^m(n+1)^n$.

[*Hint:* What can you say about the potential values for the thresholds?]

Solution:

- (c) Use these to show that the VC dimension of the class of binary-weight, threshold neural networks is $\mathcal{O}((m+n)\log(n))$.

Solution:

12. (Importance of random initialization) Consider a neural network consisting of (resp.) the input layer x , hidden layer y , hidden layer z , followed by the output node f . Suppose that all the nodes in all the layers compute a “standard” sigmoid. Also

suppose that every node in a layer is connected to every node in the next layer (*i.e.*, each layer is fully connected).

Now suppose that all the weights are initialized to 0, and suppose we start performing SGD using backprop, with a fixed learning rate. Show that at every time step, all the edge weights in a layer are equal.

Solution:

13. Let us consider networks in which each node computes a rectified linear (ReLU) function, and show how they can compute very “spiky” functions of the input variables. For this exercise, we restrict to one-variable.

Solution:

- (a) Consider a single input x . Show how to compute a “triangle wave” using one hidden layer (constant number of nodes) connected to the input, followed by one output f . Formally, we should have $f(x) = 0$ for $x \leq 0$, $f(x) = 2x$ for $0 \leq x \leq 1/2$, $f(x) = 2(1 - x)$ for $1/2 \leq x \leq 1$, and $f(x) = 0$ for $x \geq 1$.

[*Hint:* Choose the thresholds for the ReLU’s appropriately]

Solution:

- (b) What happens if you stack the network on top of itself? (Describe the function obtained).

[Formally, this means the output of the network you constructed above is fed as the input to an identical network, and we are interested in the final output function.]

Solution:

- (c) Prove that there is a ReLU network with one input variable x , $2k + \mathcal{O}(1)$ layers, all coefficients and thresholds being constants, that computes a function that has 2^k “peaks” in the interval $[0, 1]$.

[The function above can be shown to be impossible to approximate using a small depth ReLU network, without an exponential blow-up in the width.]

Solution:

14. In this exercise, we make a simple observation that width isn’t as “necessary” as depth. Consider a network in which each node computes a rectified linear (ReLU) unit – specifically the function at each node is of the form $\max\{0, a_1x_1 + a_2x_2 + \dots + a_mx_m + b\}$, for a node that has inputs $\{x_1, \dots, x_m\}$. Note that different nodes could have different coefficients and offsets (b above is called the offset).

Consider a network with one input x , connect to n nodes in a hidden layer, which are in turn connected to the output node, denoted f . Show that one can construct a depth $n + \mathcal{O}(1)$ network, with just 3 nodes in each layer, to compute the same f .

[*Hint:* Three nodes allow you to “carry over” the input; ReLU’s are important for this]

Solution: